

# Deep Learning - Musterlösung Übung 3

Convolutional Neural Networks

Fachhochschule Südwestfalen

23. Oktober 2025

## Hinweise zur Musterlösung

Diese Musterlösung enthält detaillierte mathematische Herleitungen und vollständige Implementierungen für CNNs.

## 1 Convolution-Mathematik - Lösungen

### 1.1 Aufgabe 1.1: Grundlegende Convolution

Gegeben:

$$\text{Input: } X = \begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{pmatrix} \quad (1)$$

$$\text{Kernel: } K = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix} \quad (2)$$

**Valid Convolution (ohne Padding):**

Output-Größe:  $(3 - 2 + 1) \times (3 - 2 + 1) = 2 \times 2$

$$Y[0, 0] = \sum_{i=0}^1 \sum_{j=0}^1 X[i, j] \cdot K[i, j] \quad (3)$$

$$= X[0, 0] \cdot K[0, 0] + X[0, 1] \cdot K[0, 1] + X[1, 0] \cdot K[1, 0] + X[1, 1] \cdot K[1, 1] \quad (4)$$

$$= 1 \cdot 1 + 2 \cdot 0 + 4 \cdot 0 + 5 \cdot (-1) = 1 - 5 = -4 \quad (5)$$

$$Y[0, 1] = X[0, 1] \cdot 1 + X[0, 2] \cdot 0 + X[1, 1] \cdot 0 + X[1, 2] \cdot (-1) \quad (6)$$

$$= 2 \cdot 1 + 3 \cdot 0 + 5 \cdot 0 + 6 \cdot (-1) = 2 - 6 = -4 \quad (7)$$

$$Y[1, 0] = X[1, 0] \cdot 1 + X[1, 1] \cdot 0 + X[2, 0] \cdot 0 + X[2, 1] \cdot (-1) \quad (8)$$

$$= 4 \cdot 1 + 5 \cdot 0 + 7 \cdot 0 + 8 \cdot (-1) = 4 - 8 = -4 \quad (9)$$

$$Y[1, 1] = X[1, 1] \cdot 1 + X[1, 2] \cdot 0 + X[2, 1] \cdot 0 + X[2, 2] \cdot (-1) \quad (10)$$

$$= 5 \cdot 1 + 6 \cdot 0 + 8 \cdot 0 + 9 \cdot (-1) = 5 - 9 = -4 \quad (11)$$

$$Y = \begin{pmatrix} -4 & -4 \\ -4 & -4 \end{pmatrix}$$

**Same Convolution (mit Padding):**

Padding =  $\lfloor \frac{2}{2} \rfloor = 1$ , erweiterte Matrix:

$$X_{padded} = \begin{pmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 3 & 0 \\ 0 & 4 & 5 & 6 & 0 \\ 0 & 7 & 8 & 9 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (12)$$

Output-Größe:  $3 \times 3$  (gleich wie Input)

$$Y_{same} = \begin{pmatrix} 1 & 2 & 3 \\ 4 & -4 & -4 \\ 7 & -4 & -4 \end{pmatrix} \quad (13)$$

## 1.2 Aufgabe 1.2: Multi-Channel Convolution

**Gegeben:** 3-Kanal Input ( $3 \times 3 \times 3$ ), 2 Filter ( $2 \times 2 \times 3$ )

**Mathematische Formulierung:**

$$Y^{(f)}[i, j] = \sum_{c=0}^{C-1} \sum_{u=0}^{K-1} \sum_{v=0}^{K-1} X^{(c)}[i+u, j+v] \cdot W^{(f,c)}[u, v] + b^{(f)} \quad (14)$$

Für Filter 1:

$$Y^{(1)}[0, 0] = \sum_{c=0}^2 \sum_{u=0}^1 \sum_{v=0}^1 X^{(c)}[u, v] \cdot W^{(1,c)}[u, v] + b^{(1)} \quad (15)$$

**Implementierung:**

```

1 import numpy as np
2
3 def convolution_3d(input_volume, filters, biases, stride=1, padding
4 =0):
5     """
6     3D Convolution for multi-channel inputs
7
8     Args:
9         input_volume: Shape (H, W, C)
10        filters: Shape (F, K, K, C) - F filters, each KxK with C
11        channels
12        biases: Shape (F,)
13        stride: Stride

```

```

12         padding: Padding
13
14     Returns:
15         output: Shape (H_out, W_out, F)
16     """
17     H, W, C = input_volume.shape
18     F, K, _, _ = filters.shape
19
20     # Add padding
21     if padding > 0:
22         input_padded = np.pad(input_volume,
23                                ((padding, padding), (padding, padding),
24                                 (0, 0)),
25                                mode='constant')
26     else:
27         input_padded = input_volume
28
29     # Calculate output dimensions
30     H_out = (H + 2*padding - K) // stride + 1
31     W_out = (W + 2*padding - K) // stride + 1
32
33     # Initialize output
34     output = np.zeros((H_out, W_out, F))
35
36     # Perform convolution
37     for f in range(F): # For each filter
38         for i in range(H_out):
39             for j in range(W_out):
40                 # Extract region
41                 region = input_padded[i*stride:i*stride+K,
42                                       j*stride:j*stride+K, :]
43
44                 # Convolution operation
45                 output[i, j, f] = np.sum(region * filters[f]) +
46                                     biases[f]
47
48     return output
49
50 # Example usage
51 input_vol = np.random.randn(5, 5, 3)
52 filters = np.random.randn(8, 3, 3, 3) # 8 filters, 3x3, 3 channels
53 biases = np.random.randn(8)
54
55 result = convolution_3d(input_vol, filters, biases, padding=1)
56 print(f"Output shape: {result.shape}")

```

### 1.3 Aufgabe 1.3: Pooling-Operationen

Max Pooling ( $2 \times 2$ ):

$$\text{Input: } X = \begin{pmatrix} 1 & 3 & 2 & 4 \\ 5 & 6 & 1 & 2 \\ 3 & 2 & 4 & 7 \\ 1 & 8 & 3 & 9 \end{pmatrix} \quad (16)$$

$$Y_{max}[0,0] = \max(1, 3, 5, 6) = 6 \quad (17)$$

$$Y_{max}[0,1] = \max(2, 4, 1, 2) = 4 \quad (18)$$

$$Y_{max}[1,0] = \max(3, 2, 1, 8) = 8 \quad (19)$$

$$Y_{max}[1,1] = \max(4, 7, 3, 9) = 9 \quad (20)$$

$$Y_{max} = \begin{pmatrix} 6 & 4 \\ 8 & 9 \end{pmatrix}$$

**Average Pooling:**

$$Y_{avg}[0,0] = \frac{1 + 3 + 5 + 6}{4} = 3.75 \quad (21)$$

$$Y_{avg}[0,1] = \frac{2 + 4 + 1 + 2}{4} = 2.25 \quad (22)$$

$$Y_{avg}[1,0] = \frac{3 + 2 + 1 + 8}{4} = 3.5 \quad (23)$$

$$Y_{avg}[1,1] = \frac{4 + 7 + 3 + 9}{4} = 5.75 \quad (24)$$

$$Y_{avg} = \begin{pmatrix} 3.75 & 2.25 \\ 3.5 & 5.75 \end{pmatrix}$$

## 2 CNN-Implementierung - Musterlösung

### 2.1 Aufgabe 2.1: CNN von Grund auf

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 class Conv2D:
5     def __init__(self, in_channels, out_channels, kernel_size, stride
6         =1, padding=0):
7         self.in_channels = in_channels
8         self.out_channels = out_channels
9         self.kernel_size = kernel_size
10        self.stride = stride
11        self.padding = padding
12
13        # He initialization for weights
14        self.weights = np.random.randn(out_channels, in_channels,
```

```
14         kernel_size, kernel_size) * np.
15             sqrt(2.0 / (in_channels *
16                 kernel_size * kernel_size))
17
18     self.biases = np.zeros(out_channels)
19
20     # For backpropagation
21     self.last_input = None
22     self.dW = None
23     self.db = None
24
25     def forward(self, x):
26         """Forward pass"""
27         self.last_input = x
28         batch_size, in_channels, height, width = x.shape
29
30         # Add padding
31         if self.padding > 0:
32             x_padded = np.pad(x, ((0, 0), (0, 0),
33                 (self.padding, self.padding),
34                 (self.padding, self.padding)),
35                 mode='constant')
36         else:
37             x_padded = x
38
39         # Calculate output dimensions
40         out_height = (height + 2*self.padding - self.kernel_size) //
41             self.stride + 1
42         out_width = (width + 2*self.padding - self.kernel_size) //
43             self.stride + 1
44
45         # Initialize output
46         output = np.zeros((batch_size, self.out_channels, out_height,
47             out_width))
48
49         # Convolution operation
50         for b in range(batch_size):
51             for f in range(self.out_channels):
52                 for i in range(out_height):
53                     for j in range(out_width):
54                         # Extract region
55                         region = x_padded[b, :,
56                             i*self.stride:i*self.stride+
57                             self.kernel_size,
58                             j*self.stride:j*self.stride+
59                             self.kernel_size]
60
61                         # Convolution
62                         output[b, f, i, j] = np.sum(region * self.
63                             weights[f]) + self.biases[f]
64
65     return output
```

```

57
58 def backward(self, dout):
59     """Backward pass"""
60     batch_size, in_channels, height, width = self.last_input.
        shape
61     _, out_channels, out_height, out_width = dout.shape
62
63     # Add padding to input
64     if self.padding > 0:
65         x_padded = np.pad(self.last_input, ((0, 0), (0, 0),
66                                             (self.padding, self.
67                                             padding),
68                                             (self.padding, self.
69                                             padding))),
70                               mode='constant')
71     else:
72         x_padded = self.last_input
73
74     # Initialize gradients
75     self.dW = np.zeros_like(self.weights)
76     self.db = np.zeros_like(self.biases)
77     dx_padded = np.zeros_like(x_padded)
78
79     # Compute gradients
80     for b in range(batch_size):
81         for f in range(out_channels):
82             for i in range(out_height):
83                 for j in range(out_width):
84                     # Gradient w.r.t. weights
85                     region = x_padded[b, :,
86                                     i*self.stride:i*self.stride+
87                                     self.kernel_size,
88                                     j*self.stride:j*self.stride+
89                                     self.kernel_size]
90                     self.dW[f] += dout[b, f, i, j] * region
91
92                     # Gradient w.r.t. bias
93                     self.db[f] += dout[b, f, i, j]
94
95                     # Gradient w.r.t. input
96                     dx_padded[b, :,
97                             i*self.stride:i*self.stride+self.
98                             kernel_size,
99                             j*self.stride:j*self.stride+self.
100                             kernel_size] += \
101                         dout[b, f, i, j] * self.weights[f]
102
103     # Remove padding from gradient
104     if self.padding > 0:
105         dx = dx_padded[:, :, self.padding:-self.padding, self.
106                         padding:-self.padding]

```

```
100         else:
101             dx = dx_padded
102
103         return dx
104
105 class MaxPool2D:
106     def __init__(self, pool_size=2, stride=2):
107         self.pool_size = pool_size
108         self.stride = stride
109         self.mask = None
110
111     def forward(self, x):
112         batch_size, channels, height, width = x.shape
113
114         out_height = height // self.stride
115         out_width = width // self.stride
116
117         output = np.zeros((batch_size, channels, out_height,
118                             out_width))
119         self.mask = np.zeros_like(x)
120
121         for b in range(batch_size):
122             for c in range(channels):
123                 for i in range(out_height):
124                     for j in range(out_width):
125                         # Extract pooling region
126                         region = x[b, c,
127                                   i*self.stride:i*self.stride+self.
128                                       pool_size,
129                                   j*self.stride:j*self.stride+self.
130                                       pool_size]
131
132                         # Max pooling
133                         max_val = np.max(region)
134                         output[b, c, i, j] = max_val
135
136                         # Store mask for backpropagation
137                         mask_region = (region == max_val)
138                         self.mask[b, c,
139                                   i*self.stride:i*self.stride+self.
140                                       pool_size,
141                                   j*self.stride:j*self.stride+self.
142                                       pool_size] = mask_region
143
144         return output
145
146     def backward(self, dout):
147         batch_size, channels, out_height, out_width = dout.shape
148         dx = np.zeros_like(self.mask)
149
150         for b in range(batch_size):
```

```
146         for c in range(channels):
147             for i in range(out_height):
148                 for j in range(out_width):
149                     # Distribute gradient to max element
150                     dx[b, c,
151                        i*self.stride:i*self.stride+self.pool_size,
152                        j*self.stride:j*self.stride+self.pool_size
153                        ] += \
154                         dout[b, c, i, j] * self.mask[b, c,
155                                                         i*self.stride:
156                                                         i*self.
157                                                         stride+self
158                                                         .pool_size,
159                                                         j*self.stride:
160                                                         j*self.
161                                                         stride+self
162                                                         .pool_size]
163
164         return dx
165
166 class ReLU:
167     def __init__(self):
168         self.mask = None
169
170     def forward(self, x):
171         self.mask = x > 0
172         return np.maximum(0, x)
173
174     def backward(self, dout):
175         return dout * self.mask
176
177 class Dense:
178     def __init__(self, in_features, out_features):
179         self.weights = np.random.randn(out_features, in_features) *
180             np.sqrt(2.0 / in_features)
181         self.biases = np.zeros(out_features)
182         self.last_input = None
183         self.dW = None
184         self.db = None
185
186     def forward(self, x):
187         # Flatten input if needed
188         if len(x.shape) > 2:
189             x = x.reshape(x.shape[0], -1)
190
191         self.last_input = x
192         return x @ self.weights.T + self.biases
193
194     def backward(self, dout):
195         self.dW = dout.T @ self.last_input
```



```
188         self.db = np.sum(dout, axis=0)
189         dx = dout @ self.weights
190
191         # Reshape if needed
192         if hasattr(self, 'input_shape'):
193             dx = dx.reshape(self.input_shape)
194
195         return dx
196
197 class SimpleCNN:
198     def __init__(self):
199         # Architecture: Conv -> ReLU -> MaxPool -> Conv -> ReLU ->
200             MaxPool -> Dense -> Softmax
201         self.conv1 = Conv2D(1, 6, 5, padding=2) # 28x28x1 -> 28x28x6
202         self.relu1 = ReLU()
203         self.pool1 = MaxPool2D(2, 2) # 28x28x6 -> 14x14x6
204
205         self.conv2 = Conv2D(6, 16, 5) # 14x14x6 -> 10x10x16
206         self.relu2 = ReLU()
207         self.pool2 = MaxPool2D(2, 2) # 10x10x16 -> 5x5x16
208
209         self.dense1 = Dense(5*5*16, 120)
210         self.relu3 = ReLU()
211         self.dense2 = Dense(120, 84)
212         self.relu4 = ReLU()
213         self.dense3 = Dense(84, 10) # 10 classes
214
215         self.layers = [self.conv1, self.relu1, self.pool1,
216             self.conv2, self.relu2, self.pool2,
217             self.dense1, self.relu3,
218             self.dense2, self.relu4, self.dense3]
219
220     def forward(self, x):
221         for layer in self.layers:
222             x = layer.forward(x)
223         return x
224
225     def backward(self, dout):
226         for layer in reversed(self.layers):
227             dout = layer.backward(dout)
228         return dout
229
230     def softmax(self, x):
231         exp_x = np.exp(x - np.max(x, axis=1, keepdims=True))
232         return exp_x / np.sum(exp_x, axis=1, keepdims=True)
233
234     def cross_entropy_loss(self, y_true, y_pred):
235         y_pred = self.softmax(y_pred)
236         batch_size = y_true.shape[0]
237         log_likelihood = -np.log(y_pred[range(batch_size), y_true])
238         return np.mean(log_likelihood)
```

```
238
239     def predict(self, x):
240         output = self.forward(x)
241         return np.argmax(self.softmax(output), axis=1)
```

## 2.2 Aufgabe 2.2: Training und Evaluierung

```
1  # Training function
2  def train_cnn(model, X_train, y_train, X_val, y_val, epochs=10,
3      batch_size=32, learning_rate=0.001):
4      train_losses = []
5      val_accuracies = []
6
7      for epoch in range(epochs):
8          epoch_loss = 0
9          num_batches = 0
10
11         # Shuffle training data
12         indices = np.random.permutation(len(X_train))
13         X_train_shuffled = X_train[indices]
14         y_train_shuffled = y_train[indices]
15
16         # Mini-batch training
17         for i in range(0, len(X_train), batch_size):
18             batch_X = X_train_shuffled[i:i+batch_size]
19             batch_y = y_train_shuffled[i:i+batch_size]
20
21             # Forward pass
22             output = model.forward(batch_X)
23
24             # Compute loss
25             loss = model.cross_entropy_loss(batch_y, output)
26             epoch_loss += loss
27             num_batches += 1
28
29             # Backward pass
30             # Gradient of cross-entropy + softmax
31             y_pred = model.softmax(output)
32             dout = y_pred.copy()
33             dout[range(len(batch_y)), batch_y] -= 1
34             dout /= len(batch_y)
35
36             model.backward(dout)
37
38             # Update weights
39             update_weights(model, learning_rate)
40
41         # Average loss
42         avg_loss = epoch_loss / num_batches
43         train_losses.append(avg_loss)
```

```
43
44     # Validation accuracy
45     val_pred = model.predict(X_val)
46     val_acc = np.mean(val_pred == y_val)
47     val_accuracies.append(val_acc)
48
49     print(f"Epoch {epoch+1}/{epochs}, Loss: {avg_loss:.4f}, Val
        Acc: {val_acc:.4f}")
50
51     return train_losses, val_accuracies
52
53 def update_weights(model, learning_rate):
54     """Update all trainable parameters"""
55     for layer in model.layers:
56         if hasattr(layer, 'weights'):
57             layer.weights -= learning_rate * layer.dW
58             layer.biases -= learning_rate * layer.db
59
60 # Synthetic data for testing
61 def create_synthetic_data(n_samples=1000):
62     """Create simple synthetic image data"""
63     X = np.random.randn(n_samples, 1, 28, 28)
64     y = np.random.randint(0, 10, n_samples)
65     return X, y
66
67 # Test the model
68 X_train, y_train = create_synthetic_data(1000)
69 X_val, y_val = create_synthetic_data(200)
70
71 model = SimpleCNN()
72 train_losses, val_accs = train_cnn(model, X_train, y_train, X_val,
    y_val,
73                                     epochs=5, batch_size=16,
    learning_rate=0.01)
74
75 # Plot results
76 plt.figure(figsize=(12, 4))
77 plt.subplot(1, 2, 1)
78 plt.plot(train_losses)
79 plt.title('Training Loss')
80 plt.xlabel('Epoch')
81 plt.ylabel('Loss')
82
83 plt.subplot(1, 2, 2)
84 plt.plot(val_accs)
85 plt.title('Validation Accuracy')
86 plt.xlabel('Epoch')
87 plt.ylabel('Accuracy')
88 plt.tight_layout()
89 plt.show()
```

### 3 Backpropagation in CNNs - Lösungen

#### 3.1 Aufgabe 3.1: Convolution Backpropagation

##### Mathematische Herleitung:

Für eine Convolution-Schicht mit Output  $Y = X * W + b$ :

##### Gradient bezüglich Weights:

$$\frac{\partial L}{\partial W[f, c, u, v]} = \sum_{i, j} \frac{\partial L}{\partial Y[f, i, j]} \cdot X[c, i \cdot s + u, j \cdot s + v] \quad (25)$$

##### Gradient bezüglich Input:

$$\frac{\partial L}{\partial X[c, i, j]} = \sum_{f, u, v} \frac{\partial L}{\partial Y[f, i', j']} \cdot W[f, c, u, v] \quad (26)$$

wobei  $i' = \frac{i-u}{s}$ ,  $j' = \frac{j-v}{s}$  (wenn ganzzahlig und im gültigen Bereich).

##### Numerisches Beispiel:

Gegeben: - Input:  $X$  ( $1 \times 1 \times 3 \times 3$ ) - Kernel:  $W$  ( $1 \times 1 \times 2 \times 2$ ) - Output:  $Y$  ( $1 \times 1 \times 2 \times 2$ ) - Gradient:  $\frac{\partial L}{\partial Y}$  ( $1 \times 1 \times 2 \times 2$ )

$$\frac{\partial L}{\partial Y} = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \quad (27)$$

##### Gradient bezüglich Kernel:

$$\frac{\partial L}{\partial W[0, 0]} = X[0, 0] \cdot 1 + X[0, 1] \cdot 2 + X[1, 0] \cdot 3 + X[1, 1] \cdot 4 \quad (28)$$

$$\frac{\partial L}{\partial W[0, 1]} = X[0, 1] \cdot 1 + X[0, 2] \cdot 2 + X[1, 1] \cdot 3 + X[1, 2] \cdot 4 \quad (29)$$

$$\frac{\partial L}{\partial W[1, 0]} = X[1, 0] \cdot 1 + X[1, 1] \cdot 2 + X[2, 0] \cdot 3 + X[2, 1] \cdot 4 \quad (30)$$

$$\frac{\partial L}{\partial W[1, 1]} = X[1, 1] \cdot 1 + X[1, 2] \cdot 2 + X[2, 1] \cdot 3 + X[2, 2] \cdot 4 \quad (31)$$

#### 3.2 Aufgabe 3.2: Pooling Backpropagation

##### Max Pooling Gradient:

Das Gradient wird nur an die Position weitergegeben, die das Maximum hatte:

$$\frac{\partial L}{\partial X[i, j]} = \begin{cases} \frac{\partial L}{\partial Y[i', j']} & \text{wenn } X[i, j] = \max(\text{pooling region}) \\ 0 & \text{sonst} \end{cases} \quad (32)$$

##### Implementierung:

```
1 def max_pool_backward_detailed(dout, x, pool_size=2, stride=2):
2     """
3     Detailed implementation of max pooling backward pass
4     """
```

```
5     batch_size, channels, height, width = x.shape
6     out_height, out_width = dout.shape[2], dout.shape[3]
7
8     dx = np.zeros_like(x)
9
10    for b in range(batch_size):
11        for c in range(channels):
12            for i in range(out_height):
13                for j in range(out_width):
14                    # Get pooling region
15                    h_start = i * stride
16                    h_end = h_start + pool_size
17                    w_start = j * stride
18                    w_end = w_start + pool_size
19
20                    region = x[b, c, h_start:h_end, w_start:w_end]
21
22                    # Find position of maximum
23                    max_pos = np.unravel_index(np.argmax(region),
24                                              region.shape)
25
26                    # Pass gradient to max position
27                    dx[b, c, h_start + max_pos[0], w_start + max_pos
28                      [1]] += dout[b, c, i, j]
29
30    return dx
```

## 4 Advanced Topics - Lösungen

### 4.1 Aufgabe 4.1: Data Augmentation

```
1 def augment_data(X, y):
2     """Data augmentation for image classification"""
3     augmented_X = []
4     augmented_y = []
5
6     for i in range(len(X)):
7         image = X[i].copy()
8         label = y[i]
9
10        # Original image
11        augmented_X.append(image)
12        augmented_y.append(label)
13
14        # Horizontal flip
15        flipped = np.flip(image, axis=-1)
16        augmented_X.append(flipped)
17        augmented_y.append(label)
18
19        # Rotation (simplified - small angle)
```

```

20     angle = np.random.uniform(-15, 15)
21     rotated = rotate_image(image, angle)
22     augmented_X.append(rotated)
23     augmented_y.append(label)
24
25     # Brightness adjustment
26     bright_factor = np.random.uniform(0.8, 1.2)
27     brightened = np.clip(image * bright_factor, 0, 1)
28     augmented_X.append(brightened)
29     augmented_y.append(label)
30
31     # Noise addition
32     noise = np.random.normal(0, 0.1, image.shape)
33     noisy = np.clip(image + noise, 0, 1)
34     augmented_X.append(noisy)
35     augmented_y.append(label)
36
37     return np.array(augmented_X), np.array(augmented_y)
38
39 def rotate_image(image, angle):
40     """Simple rotation implementation"""
41     # This would typically use scipy.ndimage.rotate
42     # Here's a placeholder implementation
43     return image # Simplified for this example

```

## 4.2 Aufgabe 4.2: Transfer Learning

### Konzeptuelle Erklärung:

Transfer Learning nutzt vortrainierte Modelle: 1. **Feature Extraction:** Frostore frühe Schichten, trainiere nur Klassifikator 2. **Fine-tuning:** Trainiere alle Schichten mit kleiner Learning Rate 3. **Progressive Unfreezing:** Schrittweise Freigabe von Schichten

### Implementierung:

```

1 class TransferCNN(SimpleCNN):
2     def __init__(self, pretrained_features=None, num_classes=10):
3         super().__init__()
4
5         if pretrained_features:
6             # Load pretrained convolutional layers
7             self.conv1 = pretrained_features['conv1']
8             self.conv2 = pretrained_features['conv2']
9
10            # Freeze convolutional layers
11            self.freeze_conv_layers()
12
13            # Replace classifier
14            self.dense3 = Dense(84, num_classes)
15
16    def freeze_conv_layers(self):
17        """Freeze convolutional layers for feature extraction"""
18        self.conv1.trainable = False

```

```
19         self.conv2.trainable = False
20
21     def unfreeze_conv_layers(self):
22         """Unfreeze for fine-tuning"""
23         self.conv1.trainable = True
24         self.conv2.trainable = True
25
26     def fine_tune(self, X_train, y_train, epochs=5, learning_rate
27                 =0.0001):
28         """Fine-tuning with small learning rate"""
29         self.unfreeze_conv_layers()
30
31         # Train with very small learning rate
32         train_losses, _ = train_cnn(self, X_train, y_train, X_train,
33                                     y_train,
34                                     epochs=epochs, learning_rate=
35                                     learning_rate)
36
37     return train_losses
```

## Zusammenfassung und Best Practices

### CNN Design Principles

- **Receptive Field:** Schrittweise Vergrößerung durch mehrere kleine Kernel
- **Channel Progression:** Mehr Kanäle in tieferen Schichten
- **Pooling Strategy:** Max für Features, Average für globale Information
- **Activation Choice:** ReLU für Hidden Layers, Softmax für Klassifikation

### Training Optimizations

- **Batch Normalization:** Stabilisiert Training
- **Dropout:** Regularisierung gegen Overfitting
- **Learning Rate Scheduling:** Adaptive Anpassung
- **Early Stopping:** Verhindert Overfitting