

Deep Learning - Musterlösung Übung 5

Generative Modelle und Fortgeschrittenes Deep Learning

Fachhochschule Südwestfalen

23. Oktober 2025

Hinweise zur Musterlösung

Diese Musterlösung bietet umfassende mathematische Herleitungen und praktische Implementierungen für generative Modelle und fortgeschrittene Deep Learning Techniken.

1 Autoencoders - Lösungen

1.1 Aufgabe 1.1: Autoencoder-Mathematik

Encoder-Decoder Architektur:

$$\mathbf{z} = f_{\text{enc}}(\mathbf{x}) = \sigma(\mathbf{W}_e \mathbf{x} + \mathbf{b}_e) \quad (1)$$

$$\hat{\mathbf{x}} = f_{\text{dec}}(\mathbf{z}) = \sigma(\mathbf{W}_d \mathbf{z} + \mathbf{b}_d) \quad (2)$$

$$L(\mathbf{x}, \hat{\mathbf{x}}) = \|\mathbf{x} - \hat{\mathbf{x}}\|^2 \quad (3)$$

Dimensionsanalyse: Für MNIST-Daten ($\mathbf{x} \in \mathbb{R}^{784}$) und latenten Raum ($\mathbf{z} \in \mathbb{R}^{32}$):

$$\mathbf{W}_e \in \mathbb{R}^{32 \times 784}, \quad \mathbf{b}_e \in \mathbb{R}^{32} \quad (4)$$

$$\mathbf{W}_d \in \mathbb{R}^{784 \times 32}, \quad \mathbf{b}_d \in \mathbb{R}^{784} \quad (5)$$

Gradientenberechnung:

Decoder-Gradienten:

$$\frac{\partial L}{\partial \mathbf{W}_d} = \frac{\partial L}{\partial \hat{\mathbf{x}}} \frac{\partial \hat{\mathbf{x}}}{\partial \mathbf{W}_d} \quad (6)$$

$$= -2(\mathbf{x} - \hat{\mathbf{x}}) \odot \sigma'(\mathbf{W}_d \mathbf{z} + \mathbf{b}_d) \mathbf{z}^T \quad (7)$$

Encoder-Gradienten (Backpropagation durch Decoder):

$$\frac{\partial L}{\partial \mathbf{z}} = \mathbf{W}_d^T [-2(\mathbf{x} - \hat{\mathbf{x}}) \odot \sigma'(\mathbf{W}_d \mathbf{z} + \mathbf{b}_d)] \quad (8)$$

$$\frac{\partial L}{\partial \mathbf{W}_e} = \frac{\partial L}{\partial \mathbf{z}} \odot \sigma'(\mathbf{W}_e \mathbf{x} + \mathbf{b}_e) \mathbf{x}^T \quad (9)$$

Kompressionsverhältnis:

$$\text{Compression Ratio} = \frac{\text{Original Size}}{\text{Compressed Size}} = \frac{784}{32} = 24.5 \quad (10)$$

Vergleich mit JPEG: JPEG erreicht typischerweise 10:1 bis 50:1, aber verlustbehaftet. Der Autoencoder lernt eine datenspezifische Kompression.

PCA-Vergleich: Linearer Autoencoder mit einer Hidden Layer ist äquivalent zu PCA, wenn:

- Keine Bias-Terms verwendet werden
- MSE Loss verwendet wird
- Globales Minimum erreicht wird

Beweis: Die optimalen Gewichte \mathbf{W}_d entsprechen den ersten k Hauptkomponenten.

1.2 Aufgabe 1.2: Autoencoder-Implementierung

Standard Autoencoder:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 class StandardAutoencoder:
5     def __init__(self, input_dim, latent_dim):
6         self.input_dim = input_dim
7         self.latent_dim = latent_dim
8
9         # Xavier initialization
10        self.W_encoder = np.random.randn(latent_dim, input_dim) * np.
            sqrt(2.0 / input_dim)
11        self.b_encoder = np.zeros((latent_dim, 1))
12
13        self.W_decoder = np.random.randn(input_dim, latent_dim) * np.
            sqrt(2.0 / latent_dim)
14        self.b_decoder = np.zeros((input_dim, 1))
15
16        # For storing during forward pass
17        self.cache = {}
18
19    def sigmoid(self, x):
20        return np.where(x >= 0,
21                        1 / (1 + np.exp(-x)),
22                        np.exp(x) / (1 + np.exp(x)))
23
24    def sigmoid_derivative(self, x):
25        s = self.sigmoid(x)
26        return s * (1 - s)
27
28    def encode(self, x):
29        """Encode input to latent representation"""
30        z_pre = self.W_encoder @ x + self.b_encoder
31        z = self.sigmoid(z_pre)
32        return z, z_pre
33
34    def decode(self, z):
```

```
35     """Decode latent representation to reconstruction"""
36     x_pre = self.W_decoder @ z + self.b_decoder
37     x_reconstructed = self.sigmoid(x_pre)
38     return x_reconstructed, x_pre
39
40 def forward(self, x):
41     """Complete forward pass"""
42     # Encoder
43     z, z_pre = self.encode(x)
44
45     # Decoder
46     x_reconstructed, x_pre = self.decode(z)
47
48     # Store for backpropagation
49     self.cache = {
50         'x': x,
51         'z_pre': z_pre,
52         'z': z,
53         'x_pre': x_pre,
54         'x_reconstructed': x_reconstructed
55     }
56
57     return x_reconstructed
58
59 def backward(self, x_reconstructed, x_target):
60     """Backpropagation"""
61     # Loss gradient
62     dLoss_dx_reconstructed = 2 * (x_reconstructed - x_target)
63
64     # Decoder gradients
65     dx_pre = dLoss_dx_reconstructed * self.sigmoid_derivative(
66         self.cache['x_pre'])
67     dW_decoder = dx_pre @ self.cache['z'].T
68     db_decoder = np.sum(dx_pre, axis=1, keepdims=True)
69
70     # Encoder gradients
71     dz = self.W_decoder.T @ dx_pre
72     dz_pre = dz * self.sigmoid_derivative(self.cache['z_pre'])
73     dW_encoder = dz_pre @ self.cache['x'].T
74     db_encoder = np.sum(dz_pre, axis=1, keepdims=True)
75
76     return {
77         'dW_encoder': dW_encoder,
78         'db_encoder': db_encoder,
79         'dW_decoder': dW_decoder,
80         'db_decoder': db_decoder
81     }
82
83 def update_weights(self, gradients, learning_rate):
84     """Update weights using gradients"""
85     self.W_encoder -= learning_rate * gradients['dW_encoder']
```

```
85         self.b_encoder -= learning_rate * gradients['db_encoder']
86         self.W_decoder -= learning_rate * gradients['dW_decoder']
87         self.b_decoder -= learning_rate * gradients['db_decoder']
88
89     def train(self, X, epochs=1000, learning_rate=0.01, batch_size
90             =32):
91         """Training loop"""
92         losses = []
93
94         for epoch in range(epochs):
95             epoch_loss = 0
96             num_batches = 0
97
98             # Shuffle data
99             indices = np.random.permutation(X.shape[1])
100             X_shuffled = X[:, indices]
101
102             # Mini-batch training
103             for i in range(0, X.shape[1], batch_size):
104                 batch_X = X_shuffled[:, i:i+batch_size]
105
106                 # Forward pass
107                 x_reconstructed = self.forward(batch_X)
108
109                 # Compute loss
110                 loss = np.mean((batch_X - x_reconstructed)**2)
111                 epoch_loss += loss
112                 num_batches += 1
113
114                 # Backward pass
115                 gradients = self.backward(x_reconstructed, batch_X)
116
117                 # Update weights
118                 self.update_weights(gradients, learning_rate)
119
120             avg_loss = epoch_loss / num_batches
121             losses.append(avg_loss)
122
123             if (epoch + 1) % 100 == 0:
124                 print(f"Epoch {epoch+1}/{epochs}, Loss: {avg_loss:.6f}")
125
126         return losses
127
128     def reconstruct(self, x):
129         """Reconstruct single input"""
130         return self.forward(x.reshape(-1, 1)).flatten()
131
132 class DenoisingAutoencoder(StandardAutoencoder):
133     def __init__(self, input_dim, latent_dim, noise_factor=0.3):
134         super().__init__(input_dim, latent_dim)
```

```
134         self.noise_factor = noise_factor
135
136     def add_noise(self, x):
137         """Add Gaussian noise to input"""
138         noise = np.random.normal(0, self.noise_factor, x.shape)
139         noisy_x = x + noise
140         return np.clip(noisy_x, 0, 1) # Ensure values stay in [0,1]
141
142     def train_denoising(self, X, epochs=1000, learning_rate=0.01,
143                        batch_size=32):
144         """Training with noise"""
145         losses = []
146
147         for epoch in range(epochs):
148             epoch_loss = 0
149             num_batches = 0
150
151             # Shuffle data
152             indices = np.random.permutation(X.shape[1])
153             X_shuffled = X[:, indices]
154
155             for i in range(0, X.shape[1], batch_size):
156                 batch_X = X_shuffled[:, i:i+batch_size]
157
158                 # Add noise to input
159                 noisy_X = self.add_noise(batch_X)
160
161                 # Forward pass with noisy input
162                 x_reconstructed = self.forward(noisy_X)
163
164                 # Loss computed against clean target
165                 loss = np.mean((batch_X - x_reconstructed)**2)
166                 epoch_loss += loss
167                 num_batches += 1
168
169                 # Backward pass
170                 gradients = self.backward(x_reconstructed, batch_X)
171                 self.update_weights(gradients, learning_rate)
172
173             avg_loss = epoch_loss / num_batches
174             losses.append(avg_loss)
175
176             if (epoch + 1) % 100 == 0:
177                 print(f"Epoch {epoch+1}/{epochs}, Denoising Loss: {
178                     avg_loss:.6f}")
179
180         return losses
181
182     # Test with synthetic data
183     def create_synthetic_mnist():
184         """Create synthetic MNIST-like data"""
```

```
183     np.random.seed(42)
184
185     # Create simple patterns
186     data = []
187     for _ in range(1000):
188         # Create a 28x28 image with simple patterns
189         img = np.zeros((28, 28))
190
191         # Random rectangles, circles, lines
192         if np.random.rand() < 0.33:
193             # Rectangle
194             x1, y1 = np.random.randint(5, 15, 2)
195             x2, y2 = np.random.randint(x1+3, 25, 2)
196             img[x1:x2, y1:y2] = 1
197         elif np.random.rand() < 0.66:
198             # Circle
199             center = np.random.randint(8, 20, 2)
200             radius = np.random.randint(3, 8)
201             y, x = np.ogrid[:28, :28]
202             mask = (x - center[0])**2 + (y - center[1])**2 <= radius
203                 **2
204             img[mask] = 1
205         else:
206             # Line
207             x1, y1 = np.random.randint(0, 28, 2)
208             x2, y2 = np.random.randint(0, 28, 2)
209             # Simple line drawing
210             length = max(abs(x2-x1), abs(y2-y1))
211             for t in np.linspace(0, 1, length):
212                 x = int(x1 + t*(x2-x1))
213                 y = int(y1 + t*(y2-y1))
214                 if 0 <= x < 28 and 0 <= y < 28:
215                     img[x, y] = 1
216
217             data.append(img.flatten())
218
219     return np.array(data).T # Shape: (784, 1000)
220
221 # Example usage
222 print("Erstelle synthetische Daten...")
223 X_synthetic = create_synthetic_mnist()
224
225 print("Trainiere Standard Autoencoder...")
226 autoencoder = StandardAutoencoder(input_dim=784, latent_dim=32)
227 losses_standard = autoencoder.train(X_synthetic, epochs=500,
228     learning_rate=0.01)
229
230 print("Trainiere Denoising Autoencoder...")
231 denoising_ae = DenoisingAutoencoder(input_dim=784, latent_dim=32,
232     noise_factor=0.3)
233 losses_denoising = denoising_ae.train_denoising(X_synthetic, epochs
```

```

    =500, learning_rate=0.01)
231
232 # Visualization
233 plt.figure(figsize=(12, 4))
234
235 # Plot losses
236 plt.subplot(1, 3, 1)
237 plt.plot(losses_standard, label='Standard AE')
238 plt.plot(losses_denoising, label='Denoising AE')
239 plt.title('Training Losses')
240 plt.xlabel('Epoch')
241 plt.ylabel('MSE Loss')
242 plt.legend()
243 plt.grid(True)
244
245 # Original vs Reconstruction
246 test_idx = 0
247 original = X_synthetic[:, test_idx].reshape(28, 28)
248 reconstructed_std = autoencoder.reconstruct(X_synthetic[:, test_idx])
    .reshape(28, 28)
249
250 # Add noise for denoising test
251 noisy_input = denoising_ae.add_noise(X_synthetic[:, test_idx:test_idx
    +1])
252 reconstructed_denoising = denoising_ae.forward(noisy_input).reshape
    (28, 28)
253
254 plt.subplot(1, 3, 2)
255 plt.imshow(original, cmap='gray')
256 plt.title('Original')
257 plt.axis('off')
258
259 plt.subplot(1, 3, 3)
260 plt.imshow(reconstructed_std, cmap='gray')
261 plt.title('Rekonstruktion')
262 plt.axis('off')
263
264 plt.tight_layout()
265 plt.show()

```

2 Variational Autoencoders (VAE) - Lösungen

2.1 Aufgabe 2.1: VAE-Mathematik

VAE-Formulierung:

$$q_{\phi}(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}_{\phi}(\mathbf{x}), \boldsymbol{\sigma}_{\phi}^2(\mathbf{x})) \quad (11)$$

$$p_{\theta}(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\boldsymbol{\mu}_{\theta}(\mathbf{z}), \mathbf{I}) \quad (12)$$

$$\mathcal{L}(\mathbf{x}) = \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})}[\log p_{\theta}(\mathbf{x}|\mathbf{z})] - D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x})||p(\mathbf{z})) \quad (13)$$

KL-Divergenz geschlossen: Für $q(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\sigma}^2)$ und $p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I})$:

$$D_{KL}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z})) = \frac{1}{2} \sum_{j=1}^J (\mu_j^2 + \sigma_j^2 - \log \sigma_j^2 - 1) \quad (14)$$

Reparametrisierung-Trick:

$$\mathbf{z} = \boldsymbol{\mu} + \boldsymbol{\sigma} \odot \boldsymbol{\epsilon} \quad (15)$$

$$\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \quad (16)$$

Dies macht den Sampling-Prozess differenzierbar.

2.2 Aufgabe 2.2: VAE-Implementierung

```

1 class VariationalAutoencoder:
2     def __init__(self, input_dim, latent_dim):
3         self.input_dim = input_dim
4         self.latent_dim = latent_dim
5
6         # Encoder network (outputs mu and log_var)
7         self.W_enc_1 = np.random.randn(256, input_dim) * 0.01
8         self.b_enc_1 = np.zeros((256, 1))
9
10        self.W_mu = np.random.randn(latent_dim, 256) * 0.01
11        self.b_mu = np.zeros((latent_dim, 1))
12
13        self.W_logvar = np.random.randn(latent_dim, 256) * 0.01
14        self.b_logvar = np.zeros((latent_dim, 1))
15
16        # Decoder network
17        self.W_dec_1 = np.random.randn(256, latent_dim) * 0.01
18        self.b_dec_1 = np.zeros((256, 1))
19
20        self.W_dec_2 = np.random.randn(input_dim, 256) * 0.01
21        self.b_dec_2 = np.zeros((input_dim, 1))
22
23    def relu(self, x):
24        return np.maximum(0, x)
25
26    def relu_derivative(self, x):
27        return (x > 0).astype(float)
28
29    def sigmoid(self, x):
30        return np.where(x >= 0,
31                        1 / (1 + np.exp(-x)),
32                        np.exp(x) / (1 + np.exp(x)))
33
34    def encode(self, x):
35        """Encoder: x -> mu, log_var"""
36        h1 = self.relu(self.W_enc_1 @ x + self.b_enc_1)

```



```
37     mu = self.W_mu @ h1 + self.b_mu
38     log_var = self.W_logvar @ h1 + self.b_logvar
39     return mu, log_var, h1
40
41     def reparameterize(self, mu, log_var):
42         """Reparameterization trick"""
43         std = np.exp(0.5 * log_var)
44         eps = np.random.normal(0, 1, mu.shape)
45         z = mu + std * eps
46         return z, eps
47
48     def decode(self, z):
49         """Decoder: z -> x_reconstructed"""
50         h1 = self.relu(self.W_dec_1 @ z + self.b_dec_1)
51         x_reconstructed = self.sigmoid(self.W_dec_2 @ h1 + self.
52             b_dec_2)
53         return x_reconstructed, h1
54
55     def forward(self, x):
56         """Complete forward pass"""
57         # Encode
58         mu, log_var, h_enc = self.encode(x)
59
60         # Sample
61         z, eps = self.reparameterize(mu, log_var)
62
63         # Decode
64         x_reconstructed, h_dec = self.decode(z)
65
66         # Store for backpropagation
67         self.cache = {
68             'x': x,
69             'h_enc': h_enc,
70             'mu': mu,
71             'log_var': log_var,
72             'z': z,
73             'eps': eps,
74             'h_dec': h_dec,
75             'x_reconstructed': x_reconstructed
76         }
77
78         return x_reconstructed, mu, log_var
79
80     def compute_loss(self, x, x_reconstructed, mu, log_var):
81         """VAE loss = Reconstruction loss + KL divergence"""
82         batch_size = x.shape[1]
83
84         # Reconstruction loss (binary cross-entropy)
85         reconstruction_loss = -np.sum(
86             x * np.log(x_reconstructed + 1e-8) +
87             (1 - x) * np.log(1 - x_reconstructed + 1e-8)
```

```
87         ) / batch_size
88
89     # KL divergence
90     kl_loss = -0.5 * np.sum(1 + log_var - mu**2 - np.exp(log_var)
91                             ) / batch_size
92
93     total_loss = reconstruction_loss + kl_loss
94
95     return total_loss, reconstruction_loss, kl_loss
96
97 def generate(self, num_samples=1):
98     """Generate new samples from prior"""
99     z = np.random.normal(0, 1, (self.latent_dim, num_samples))
100     generated, _ = self.decode(z)
101     return generated
102
103 # Example VAE training (simplified)
104 def train_vae_example():
105     # Create simple synthetic data
106     X = np.random.rand(784, 100) # 100 samples of 784-dim data
107
108     vae = VariationalAutoencoder(input_dim=784, latent_dim=20)
109
110     print("Training VAE...")
111     for epoch in range(100):
112         # Forward pass
113         x_recon, mu, log_var = vae.forward(X)
114
115         # Compute loss
116         total_loss, recon_loss, kl_loss = vae.compute_loss(X, x_recon,
117                                                             mu, log_var)
118
119         if (epoch + 1) % 20 == 0:
120             print(f"Epoch {epoch+1}: Total Loss: {total_loss:.4f}, "
121                   f"Recon: {recon_loss:.4f}, KL: {kl_loss:.4f}")
122
123         # Generate new samples
124         generated = vae.generate(5)
125         print(f"Generated samples shape: {generated.shape}")
126
127 train_vae_example()
```

3 Generative Adversarial Networks (GANs) - Lösungen

3.1 Aufgabe 3.1: GAN-Theorie

Minimax-Spiel:

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z} [\log(1 - D(G(\mathbf{z})))] \quad (17)$$

Optimaler Discriminator: Für festen Generator G , der optimale Discriminator ist:

$$D^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})} \quad (18)$$

Beweis: Zu maximieren:

$$V(G, D) = \int_{\mathbf{x}} p_{\text{data}}(\mathbf{x}) \log D(\mathbf{x}) + p_g(\mathbf{x}) \log(1 - D(\mathbf{x})) d\mathbf{x} \quad (19)$$

Ableitung nach $D(\mathbf{x})$ und Nullsetzen:

$$\frac{\partial}{\partial D(\mathbf{x})} V(G, D) = \frac{p_{\text{data}}(\mathbf{x})}{D(\mathbf{x})} - \frac{p_g(\mathbf{x})}{1 - D(\mathbf{x})} = 0 \quad (20)$$

Lösung: $D^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})}$

Globales Optimum: Wenn $p_g = p_{\text{data}}$, dann $D^*(\mathbf{x}) = \frac{1}{2}$ und $V(G^*, D^*) = -\log 4$.

3.2 Aufgabe 3.2: GAN-Implementierung

```

1 class SimpleGAN:
2     def __init__(self, latent_dim=100, data_dim=784):
3         self.latent_dim = latent_dim
4         self.data_dim = data_dim
5
6         # Generator weights
7         self.G_W1 = np.random.randn(128, latent_dim) * 0.01
8         self.G_b1 = np.zeros((128, 1))
9         self.G_W2 = np.random.randn(data_dim, 128) * 0.01
10        self.G_b2 = np.zeros((data_dim, 1))
11
12        # Discriminator weights
13        self.D_W1 = np.random.randn(128, data_dim) * 0.01
14        self.D_b1 = np.zeros((128, 1))
15        self.D_W2 = np.random.randn(1, 128) * 0.01
16        self.D_b2 = np.zeros((1, 1))
17
18        def leaky_relu(self, x, alpha=0.2):
19            return np.where(x > 0, x, alpha * x)
20
21        def leaky_relu_derivative(self, x, alpha=0.2):
22            return np.where(x > 0, 1, alpha)

```

```

23
24 def sigmoid(self, x):
25     return np.where(x >= 0,
26                     1 / (1 + np.exp(-x)),
27                     np.exp(x) / (1 + np.exp(x)))
28
29 def sigmoid_derivative(self, x):
30     s = self.sigmoid(x)
31     return s * (1 - s)
32
33 def generator(self, z):
34     """Generator: z -> fake_data"""
35     h1 = self.leaky_relu(self.G_W1 @ z + self.G_b1)
36     output = self.sigmoid(self.G_W2 @ h1 + self.G_b2)
37     return output, h1
38
39 def discriminator(self, x):
40     """Discriminator: x -> probability"""
41     h1 = self.leaky_relu(self.D_W1 @ x + self.D_b1)
42     output = self.sigmoid(self.D_W2 @ h1 + self.D_b2)
43     return output, h1
44
45 def train_discriminator(self, real_data, fake_data, learning_rate
46 =0.0002):
47     """Train discriminator for one step"""
48     batch_size = real_data.shape[1]
49
50     # Forward pass on real data
51     real_output, real_h1 = self.discriminator(real_data)
52
53     # Forward pass on fake data
54     fake_output, fake_h1 = self.discriminator(fake_data)
55
56     # Discriminator loss
57     d_loss_real = -np.mean(np.log(real_output + 1e-8))
58     d_loss_fake = -np.mean(np.log(1 - fake_output + 1e-8))
59     d_loss = d_loss_real + d_loss_fake
60
61     # Gradients for real data
62     d_real_output = -1 / (real_output + 1e-8) / batch_size
63     d_real_h1_pre = d_real_output * self.sigmoid_derivative(self.
64         D_W2 @ real_h1 + self.D_b2)
65     d_D_W2_real = d_real_h1_pre @ real_h1.T
66     d_D_b2_real = np.sum(d_real_h1_pre, axis=1, keepdims=True)
67
68     d_real_h1 = self.D_W2.T @ d_real_h1_pre
69     d_real_h1_pre_2 = d_real_h1 * self.leaky_relu_derivative(self.
70         D_W1 @ real_data + self.D_b1)
71     d_D_W1_real = d_real_h1_pre_2 @ real_data.T
72     d_D_b1_real = np.sum(d_real_h1_pre_2, axis=1, keepdims=True)

```

```

71     # Gradients for fake data
72     d_fake_output = 1 / (1 - fake_output + 1e-8) / batch_size
73     d_fake_h1_pre = d_fake_output * self.sigmoid_derivative(self.
74         D_W2 @ fake_h1 + self.D_b2)
75     d_D_W2_fake = d_fake_h1_pre @ fake_h1.T
76     d_D_b2_fake = np.sum(d_fake_h1_pre, axis=1, keepdims=True)
77
78     d_fake_h1 = self.D_W2.T @ d_fake_h1_pre
79     d_fake_h1_pre_2 = d_fake_h1 * self.leaky_relu_derivative(self.
80         D_W1 @ fake_data + self.D_b1)
81     d_D_W1_fake = d_fake_h1_pre_2 @ fake_data.T
82     d_D_b1_fake = np.sum(d_fake_h1_pre_2, axis=1, keepdims=True)
83
84     # Update discriminator weights
85     self.D_W2 -= learning_rate * (d_D_W2_real + d_D_W2_fake)
86     self.D_b2 -= learning_rate * (d_D_b2_real + d_D_b2_fake)
87     self.D_W1 -= learning_rate * (d_D_W1_real + d_D_W1_fake)
88     self.D_b1 -= learning_rate * (d_D_b1_real + d_D_b1_fake)
89
90     return d_loss
91
92 def train_generator(self, z, learning_rate=0.0002):
93     """Train generator for one step"""
94     batch_size = z.shape[1]
95
96     # Generate fake data
97     fake_data, g_h1 = self.generator(z)
98
99     # Pass through discriminator
100     d_output, d_h1 = self.discriminator(fake_data)
101
102     # Generator loss (wants discriminator to output 1)
103     g_loss = -np.mean(np.log(d_output + 1e-8))
104
105     # Backpropagate through discriminator (frozen weights)
106     d_d_output = -1 / (d_output + 1e-8) / batch_size
107     d_d_h1_pre = d_d_output * self.sigmoid_derivative(self.D_W2 @
108         d_h1 + self.D_b2)
109     d_d_h1 = self.D_W2.T @ d_d_h1_pre
110     d_fake_data = self.D_W1.T @ (d_d_h1 * self.
111         leaky_relu_derivative(self.D_W1 @ fake_data + self.D_b1))
112
113     # Backpropagate through generator
114     d_g_output = d_fake_data * self.sigmoid_derivative(self.G_W2
115         @ g_h1 + self.G_b2)
116     d_G_W2 = d_g_output @ g_h1.T
117     d_G_b2 = np.sum(d_g_output, axis=1, keepdims=True)
118
119     d_g_h1 = self.G_W2.T @ d_g_output
120     d_g_h1_pre = d_g_h1 * self.leaky_relu_derivative(self.G_W1 @
121         z + self.G_b1)

```

```
116     d_G_W1 = d_g_h1_pre @ z.T
117     d_G_b1 = np.sum(d_g_h1_pre, axis=1, keepdims=True)
118
119     # Update generator weights
120     self.G_W2 -= learning_rate * d_G_W2
121     self.G_b2 -= learning_rate * d_G_b2
122     self.G_W1 -= learning_rate * d_G_W1
123     self.G_b1 -= learning_rate * d_G_b1
124
125     return g_loss
126
127 def generate_samples(self, num_samples):
128     """Generate samples from random noise"""
129     z = np.random.normal(0, 1, (self.latent_dim, num_samples))
130     generated, _ = self.generator(z)
131     return generated
132
133 # Example training
134 def train_gan_example():
135     # Synthetic real data
136     real_data = np.random.rand(784, 1000)
137
138     gan = SimpleGAN(latent_dim=100, data_dim=784)
139
140     epochs = 1000
141     batch_size = 64
142
143     print("Training GAN...")
144     for epoch in range(epochs):
145         # Random batch of real data
146         idx = np.random.randint(0, real_data.shape[1], batch_size)
147         real_batch = real_data[:, idx]
148
149         # Generate fake data
150         z = np.random.normal(0, 1, (gan.latent_dim, batch_size))
151         fake_batch, _ = gan.generator(z)
152
153         # Train discriminator
154         d_loss = gan.train_discriminator(real_batch, fake_batch)
155
156         # Train generator
157         z = np.random.normal(0, 1, (gan.latent_dim, batch_size))
158         g_loss = gan.train_generator(z)
159
160         if (epoch + 1) % 100 == 0:
161             print(f"Epoch {epoch+1}: D_loss: {d_loss:.4f}, G_loss: {g_loss:.4f}")
162
163     # Generate samples
164     samples = gan.generate_samples(10)
```

```
165     print(f"Generated {samples.shape[1]} samples of dimension {  
166           samples.shape[0]}")  
167 train_gan_example()
```

4 Vertiefende Aufgaben - Lösungen

4.1 Aufgabe 4.1: Data Augmentation

```
1 class DataAugmentation:  
2     def __init__(self):  
3         pass  
4  
5     def horizontal_flip(self, image):  
6         """Horizontal flip"""  
7         return np.fliplr(image)  
8  
9     def vertical_flip(self, image):  
10        """Vertical flip"""  
11        return np.flipud(image)  
12  
13    def rotation(self, image, angle):  
14        """Simple rotation (simplified implementation)"""  
15        # In practice, use scipy.ndimage.rotate or cv2.rotate  
16        # This is a placeholder  
17        return image  
18  
19    def gaussian_noise(self, image, mean=0, std=0.1):  
20        """Add Gaussian noise"""  
21        noise = np.random.normal(mean, std, image.shape)  
22        noisy_image = image + noise  
23        return np.clip(noisy_image, 0, 1)  
24  
25    def brightness_adjustment(self, image, factor):  
26        """Adjust brightness"""  
27        bright_image = image * factor  
28        return np.clip(bright_image, 0, 1)  
29  
30    def contrast_adjustment(self, image, factor):  
31        """Adjust contrast"""  
32        mean = np.mean(image)  
33        contrast_image = (image - mean) * factor + mean  
34        return np.clip(contrast_image, 0, 1)  
35  
36    def random_crop(self, image, crop_size):  
37        """Random crop"""  
38        h, w = image.shape  
39        ch, cw = crop_size  
40  
41        if h < ch or w < cw:
```

```

42         return image
43
44     x = np.random.randint(0, h - ch + 1)
45     y = np.random.randint(0, w - cw + 1)
46
47     return image[x:x+ch, y:y+cw]
48
49     def augment_batch(self, images, augmentation_prob=0.5):
50         """Apply random augmentations to a batch"""
51         augmented = []
52
53         for img in images:
54             # Reshape if needed
55             if img.ndim == 1:
56                 img = img.reshape(28, 28) # Assume MNIST-like
57
58             # Apply random augmentations
59             if np.random.rand() < augmentation_prob:
60                 # Choose random augmentation
61                 aug_type = np.random.choice(['flip', 'noise', '
62                     brightness', 'contrast'])
63
64                 if aug_type == 'flip':
65                     if np.random.rand() < 0.5:
66                         img = self.horizontal_flip(img)
67                     else:
68                         img = self.vertical_flip(img)
69                 elif aug_type == 'noise':
70                     img = self.gaussian_noise(img, std=0.1)
71                 elif aug_type == 'brightness':
72                     factor = np.random.uniform(0.8, 1.2)
73                     img = self.brightness_adjustment(img, factor)
74                 elif aug_type == 'contrast':
75                     factor = np.random.uniform(0.8, 1.2)
76                     img = self.contrast_adjustment(img, factor)
77
78             augmented.append(img.flatten())
79
80         return np.array(augmented)
81
82     # Example usage
83     augmenter = DataAugmentation()
84     sample_images = np.random.rand(10, 784) # 10 samples
85     augmented = augmenter.augment_batch(sample_images)
86     print(f"Augmented {len(augmented)} images")

```

4.2 Aufgabe 4.2: Gradient Clipping

```

1 class GradientClipper:
2     def __init__(self, max_norm=5.0):

```



```
3         self.max_norm = max_norm
4
5     def clip_gradients(self, gradients):
6         """Clip gradients by global norm"""
7         # Calculate global norm
8         total_norm = 0
9         for grad in gradients.values():
10             if isinstance(grad, np.ndarray):
11                 total_norm += np.sum(grad**2)
12
13         total_norm = np.sqrt(total_norm)
14
15         # Clip if necessary
16         if total_norm > self.max_norm:
17             clip_ratio = self.max_norm / total_norm
18             clipped_gradients = {}
19             for key, grad in gradients.items():
20                 if isinstance(grad, np.ndarray):
21                     clipped_gradients[key] = grad * clip_ratio
22                 else:
23                     clipped_gradients[key] = grad
24             return clipped_gradients, total_norm
25
26         return gradients, total_norm
27
28     def clip_gradients_by_value(self, gradients, min_val=-1.0,
29                                max_val=1.0):
30         """Clip gradients by value"""
31         clipped_gradients = {}
32         for key, grad in gradients.items():
33             if isinstance(grad, np.ndarray):
34                 clipped_gradients[key] = np.clip(grad, min_val,
35                                                    max_val)
36             else:
37                 clipped_gradients[key] = grad
38         return clipped_gradients
39
40 # Example usage with RNN
41 class RNNWithClipping:
42     def __init__(self, input_size, hidden_size, output_size):
43         self.W_xh = np.random.randn(hidden_size, input_size) * 0.01
44         self.W_hh = np.random.randn(hidden_size, hidden_size) * 0.01
45         self.W_hy = np.random.randn(output_size, hidden_size) * 0.01
46         self.b_h = np.zeros((hidden_size, 1))
47         self.b_y = np.zeros((output_size, 1))
48
49         self.clipper = GradientClipper(max_norm=5.0)
50
51     def forward(self, inputs):
52         """Forward pass through RNN"""
53         h = np.zeros((self.W_hh.shape[0], 1))
```

```

52     outputs = []
53     self.cache = {'inputs': inputs, 'hiddens': [h.copy()]}
54
55     for x in inputs:
56         x = x.reshape(-1, 1)
57         h = np.tanh(self.W_xh @ x + self.W_hh @ h + self.b_h)
58         y = self.W_hy @ h + self.b_y
59         outputs.append(y)
60         self.cache['hiddens'].append(h.copy())
61
62     return outputs
63
64     def backward(self, doutputs):
65         """Backward pass with gradient computation"""
66         # Simplified backward pass
67         gradients = {
68             'W_xh': np.zeros_like(self.W_xh),
69             'W_hh': np.zeros_like(self.W_hh),
70             'W_hy': np.zeros_like(self.W_hy),
71             'b_h': np.zeros_like(self.b_h),
72             'b_y': np.zeros_like(self.b_y)
73         }
74
75         # Accumulate gradients (simplified)
76         for i, dout in enumerate(doutputs):
77             gradients['W_hy'] += dout @ self.cache['hiddens'][i+1].T
78             gradients['b_y'] += dout
79             # ... (weitere Gradienten-Berechnungen)
80
81         # Clip gradients
82         clipped_gradients, grad_norm = self.clipper.clip_gradients(
            gradients)
83
84         return clipped_gradients, grad_norm
85
86     def update_weights(self, gradients, learning_rate):
87         """Update weights with clipped gradients"""
88         self.W_xh -= learning_rate * gradients['W_xh']
89         self.W_hh -= learning_rate * gradients['W_hh']
90         self.W_hy -= learning_rate * gradients['W_hy']
91         self.b_h -= learning_rate * gradients['b_h']
92         self.b_y -= learning_rate * gradients['b_y']
93
94     print("Gradient Clipping implementiert")

```

Zusammenfassung und Best Practices

Generative Modelle - Vergleich

- ****Autoencoders:**** Deterministische Kompression, gut für Dimensionsreduktion

- **VAEs:** Probabilistische Generierung, interpretierbare latente Räume
- **GANs:** Hochqualitative Samples, aber Training instabil

Training-Stabilität

- **Data Augmentation:** Erhöht Generalisierung und Robustheit
- **Gradient Clipping:** Verhindert explodierenden Gradienten
- **Proper Initialization:** Xavier/He für stabile Aktivierungen
- **Learning Rate Scheduling:** Adaptive Anpassung während Training

Praktische Tipps

- **VAE Training:** Balance zwischen Rekonstruktion und KL-Loss
- **GAN Training:** Discriminator nicht zu stark trainieren
- **Monitoring:** Visualisierung von generierten Samples
- **Evaluation:** FID, IS für quantitative Bewertung